

Parallel Progrmg: Sci & Engrg

Machine Problem 2

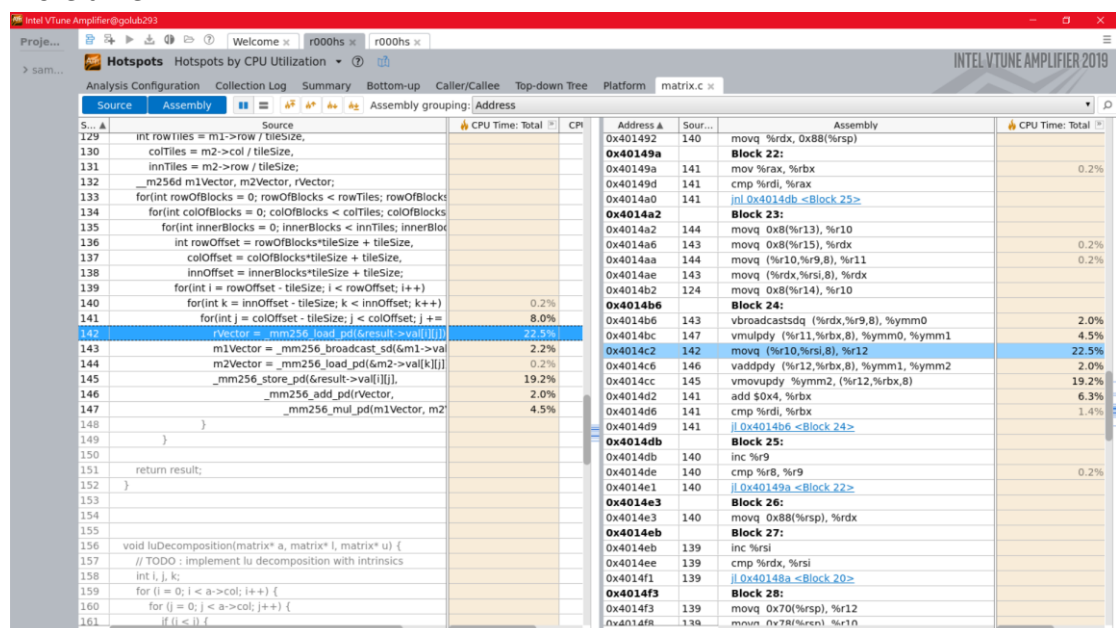
yhyuan2

- A. Q: For the version of matmul without vector intrinsics, did the compiler generate packed vector instructions? If so, pick one of the packed vector instructions it generated and briefly explain what it does. Hint: Consult the assembly tab.

A: Compiler did generate packed vector instructions. Take `vmulpd %ymm15, %ymm10, %ymm10` for example. It multiply `%ymm10` with itself and put the result into `%ymm15`.

- B. Q: Which of the vector intrinsics in your hand-vectorized program dominated your program's execution time? Briefly explain why and include a screenshot of your hand-vectorized program's top hotspot like Figure 2 in your report.

A: `rVector = _mm256_load_pd(&result->val[i][j])` dominates the execution time because it accesses result value from memory. Accessing values from memory will need to retrieve data and place them in cache to execute. So it will spend more time.



C. Q: How did the performance of your hand-vectorized version compare to the original version? Include the MFLOPs rates for both versions and briefly comment about whether or not you think the compiler's auto-vectorization was adequate and why.

A: Hand-vectorized version is much faster than original version. MFLOPs rates are listed below. I think auto-vectorization isn't adequate. Auto-vectorization is still not smart enough to find how to vectorize our code will make it faster. Also, multiplication in original version is simply row times column. It's a little bit different from the implementation of multiplication in hand-vectorized version. That's hard for compiler to vectorize original version.

	Hand-vectorized Version	Original Version
MFLOPs	5534.283367	2633.243824